

# Coherence-Based Alignment: A Preliminary Empirical Test in a Reward-Loop Gridworld

Technical Note - November 2025

Author: Abdulaziz Abdi

## Abstract

This note presents a preliminary empirical test of Coherence-Based Alignment (CBA) in a  $10 \times 10$  reward-loop gridworld. CBA introduces small penalties for incoherent action selection and loop re-entry during learning. Compared to baseline Q-learning across five seeds, CBA reduces reward-loop entrapment by  $\sim 66\%$  while preserving near-perfect goal achievement. These results provide early evidence that coherence regularization may reduce simple reward-hacking behaviors.

## 1. Overview

This technical note documents the first empirical test of *Coherence-Based Alignment (CBA)* in a controlled reinforcement learning environment. The objective is to evaluate whether a simple form of coherence regularization can reduce reward-loop entrapment—one of the most basic forms of reward hacking.

Using a  $10 \times 10$  gridworld with a deliberately designed “local reward attractor,” we compare baseline Q-learning with a modified update rule that includes a small penalty for incoherent action selection and for entering known local reward loops.

This is an early-stage *proof-of-concept* experiment. It does not constitute a full RL safety method, but it provides initial evidence that coherence-regularized learning can reduce pathological policy behavior without impairing task performance.

## 2. Background and Motivation

Many RL agents exhibit *reward-loop entrapment*: the agent converges to a policy that exploits small, repeatable rewards instead of pursuing the true sparse reward. This is one of the simplest forms of reward hacking.

*Coherence-Based Alignment (CBA)* proposes that, aside from reward maximization, policies should maintain *internal consistency*:

- actions should generally follow what the agent believes is its own best estimate,
- and repeated exploitation of narrow loops should be discouraged.

This experiment tests a minimal version of CBA as an *instantaneous regularizer*, not as a new objective function.

## 3. Environment

We use a custom  $10 \times 10$  gridworld with:

**Start:**

(0, 0)

**Goal:**

(9, 9) — reward = +100, episode terminates

**Local Reward Loop:**

A  $2 \times 2$  region in the center:

- (4, 4)
- (4, 5)
- (5, 4)
- (5, 5)

Each entry gives a small +0.5 reward that can be repeated indefinitely.  
This creates a strong local attractor.

**Step reward:**

-0.1

**Actions:**

Up, Right, Down, Left

**Transition dynamics:**

Deterministic.

The environment is designed so that naive Q-learners often *prefer the loop* instead of going to the goal.

## 4. Methods

### 4.1 Baseline Algorithm

Standard tabular *Q-learning*:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Hyperparameters:

- $\alpha = 0.1$
- $\gamma = 0.95$
- $\epsilon$ -greedy exploration with exponential decay ( $1.0 \rightarrow 0.01$  over 5000 episodes)
- Max episode length: 1000 steps
- 5000 episodes per run
- 5 random seeds

## 4.2 CBA-Regularized Update

CBA adds **two instantaneous penalties**:

1. **Loop penalty ( $L_t$ ):**  
-1 if the agent enters the loop region, else 0
2. **Incoherence penalty ( $I_t$ ):**  
-1 if the action is non-greedy under current  $Q(s)$ , else 0

Regularized TD update:

$$TD_{error} = r + \gamma \max_{a'} Q(s', a') - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [TD_{error} + \lambda(L_t + I_t)]$$

$\lambda = 0.5$  in this experiment.

## 5. Metrics

We track:

1. **Goal hits:** how often the agent reaches the true goal
2. **Average episode return**
3. **Fraction of steps spent in the loop region**

The third metric is the key measure of reward-loop entrapment.

All metrics are averaged over the **final 1000 episodes** for stability, then aggregated across the 5 seeds.

## 6. Results

**Final Aggregated Performance (Average over 5 seeds)**

Metric	Baseline ( $\lambda=0$ )	CBA ( $\lambda=0.5$ )
Average Return (last 1000 episodes)	99.79	98.79
Goal-Reaching Rate (per 100 episodes)	99.85	99.88
Fraction of Steps Spent in Loop (Loop-Time)	0.166	0.056

### Interpretation

- Both agents learn to reach the goal at the same high level of reliability.
- Baseline agent spends **16.6%** of its time inside the reward loop.
- CBA-regularized agent spends **5.6%** of its time inside the loop.
- This is an **absolute reduction of 11.0%**, or a **~66% relative reduction**.

### Conclusion

The CBA agent *substantially reduces reward-loop entrapment* without harming goal performance.

This is the intended effect: avoiding narrow local attractors while still maximizing reward.

## 7. Limitations

This is a simple, deterministic environment with tabular Q-learning.

The method has not yet been tested on:

- stochastic transitions
- larger or continuous spaces
- neural network function approximators
- multi-loop environments
- partial observability
- multi-agent settings

These tests are planned for future work.

## 8. Proposed Next Steps

1. Share this note + code with RL engineers to obtain feedback on which next environment is most meaningful.
2. Test under *stochastic transitions*.

3. Add *multiple loops* with varying reward structures.
4. Replace tabular Q with a small *DQN* to test whether CBA scales to function approximation.
5. Evaluate on environments with *misleading local optima*, such as:
  - o mountain car with shaped traps
  - o modified cliff-walking
  - o gridworlds with partial observability
6. Package results into a short arXiv preprint after 3–5 experiments.

This staged approach mirrors standard safety research methodology and avoids premature conclusions.

## 9. Repository Access

The full code used to run the experiment will be included in the accompanying GitHub repository, along with:

- complete source files
- instructions for reproduction
- seed-averaged results
- experiment logs
- environment diagrams

## 10. Summary

This preliminary test offers the first empirical evidence supporting the claim that *coherence penalties* can reduce pathological reward-seeking behavior in RL agents.

While extremely early, the results justify further investigation and more rigorous evaluations.

CBA is not a complete alignment method, but this experiment demonstrates that coherence-regularization behaves as intended in a simple RL setting. Further evaluation in richer environments is required.

## Code Repository

[https://github.com/abdulazizmohamed-dotcom/cba-gridworld-experiment/blob/main/cba\\_experiment.py](https://github.com/abdulazizmohamed-dotcom/cba-gridworld-experiment/blob/main/cba_experiment.py)